

# Resource intensive jobs: when AWS Lambda performs too well



# This is THN: Growth for the direct channel of hotels



An ecosystem of  
growth tools for hotels

We are a distributed platform for the  
hotel industry

15,000 hotels

Several million daily visits

Over 8PB of data processed monthly



# Intel (my team)

- We **enrich data** (real time or not) a lot. In order to provide value, we want the richest data possible. This is not always possible with the data we capture from the hotel's website.
- What **type of data** do we enrich? Room names & description, prices, availability, available discounts...
- We are building a **highly concurrent pipeline** to grab, process and normalize data, on the top of an event-based set of microservices written in Go and Node.JS backed by AWS.



# One need: taking photos of disparities

- Sometimes prices in sources and aggregators do not follow the hotelier guidelines
- In that case we take a photo and save it for later hotelier usage

Your search: 2 adults, 1 night [Edit](#)

<b>Check-in</b> <b>Tue 27 Dec</b> From 14:00	<b>Check-out</b> <b>Wed 28 Dec</b> Until 12:00
--	--

3 results

**Double or Twin Room (1-2 Adults)** 

Price for up to:  

Beds: 2 single beds  

20 m<sup>2</sup>  Air conditioning  Ensuite bathroom  
 Soundproofing  Free WiFi

✓ Free cancellation before 24 December 2022  
✓ NO PREPAYMENT NEEDED - pay at the property

Price for **1 night**: **US\$85**  Only 4 left on our site!

Includes taxes and charges [Reserve](#)



# High-level design

1. Another microservice captures the disparity and **produces** a message with all the information needed.
2. The **consumer** reads the event and sends it to the internal **processor**.
3. The processor launches a **browser** and navigates to the original site.
4. The processor does, if needed, some **tweaks** into the DOM.
5. The processor takes a screenshot and **uploads** it into a remote disk.

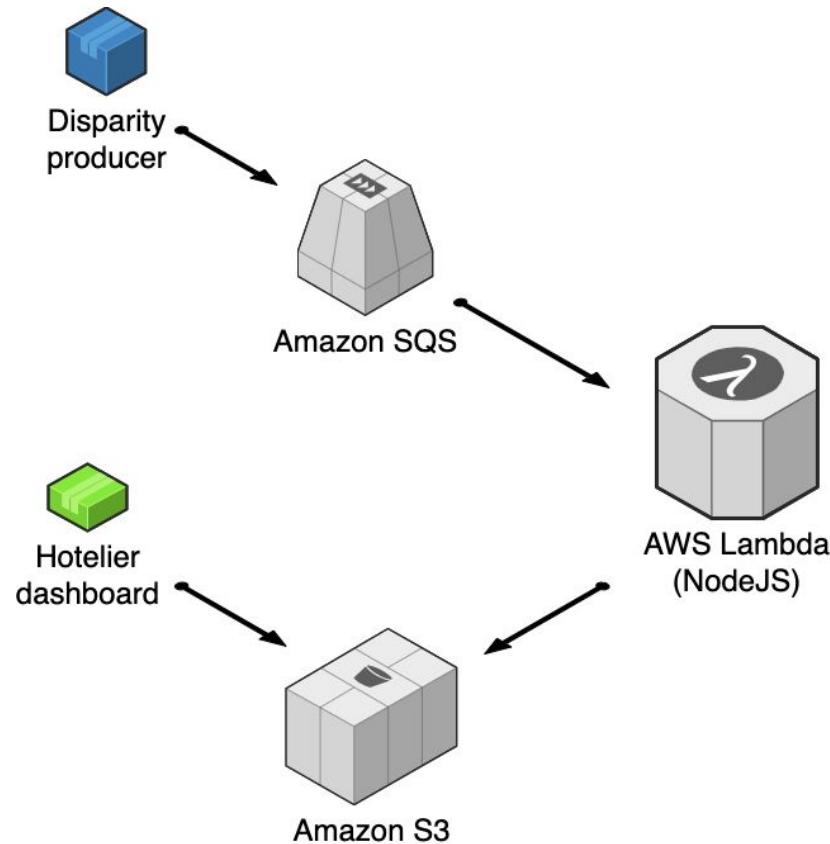


# Numbers

- >30 sources to take screenshots
- ~300K RPD → 250 RPM
- ~16s to process one screenshot
- ~6s delay to process screenshots



# Initial approach: AWS lambda to rescue



# AWS Lambda

FaaS:  
serverless  
computing

Well  
integrated in  
AWS  
ecosystem

Multiple  
SDKs,  
including  
NodeJS

Pricing:  
per allocation,  
per usage



# Problems found

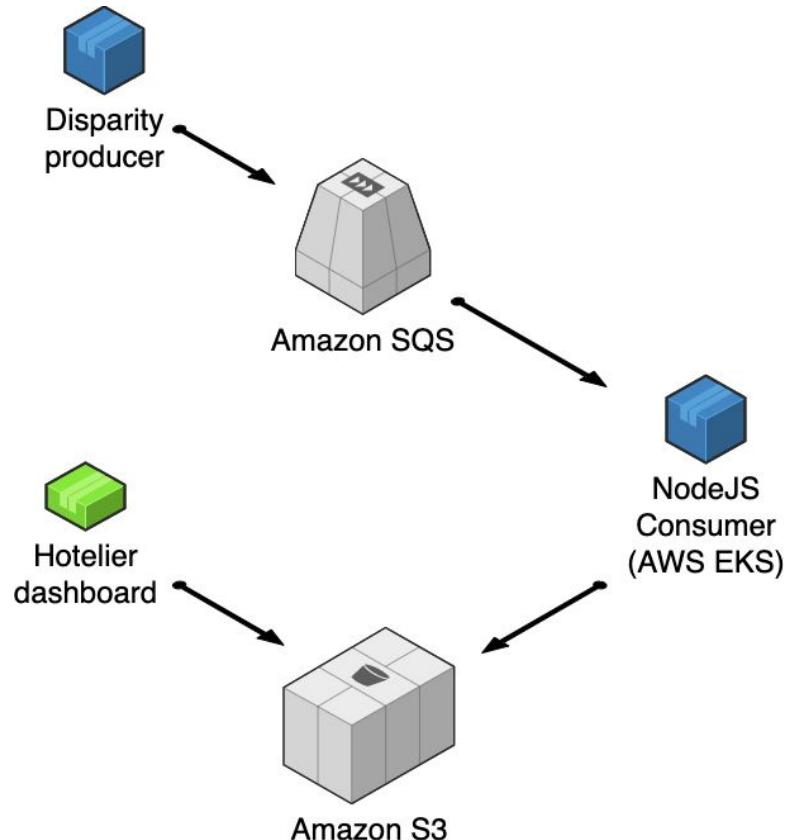


# Analysis

- Tons of requests
- High resources usage:
  - Time running
  - Memory
  - CPU
- Retries



# Reengineering with AWS EKS



# Problems found



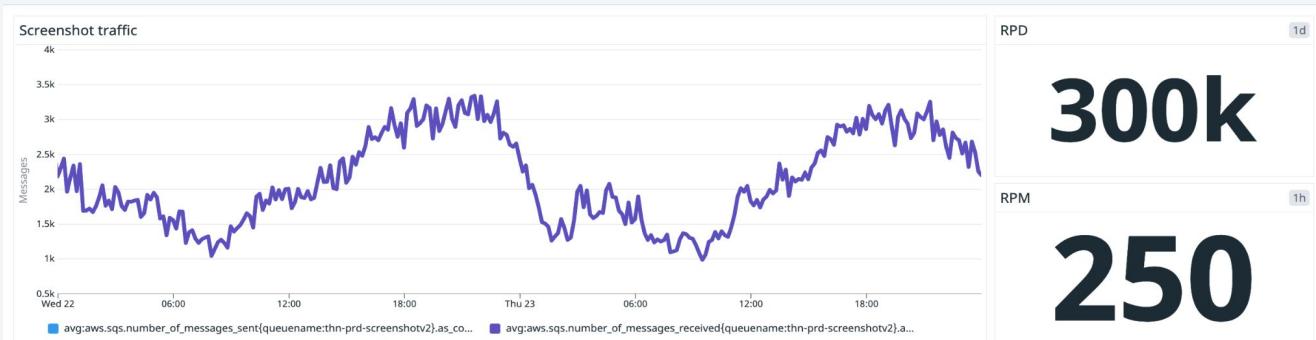
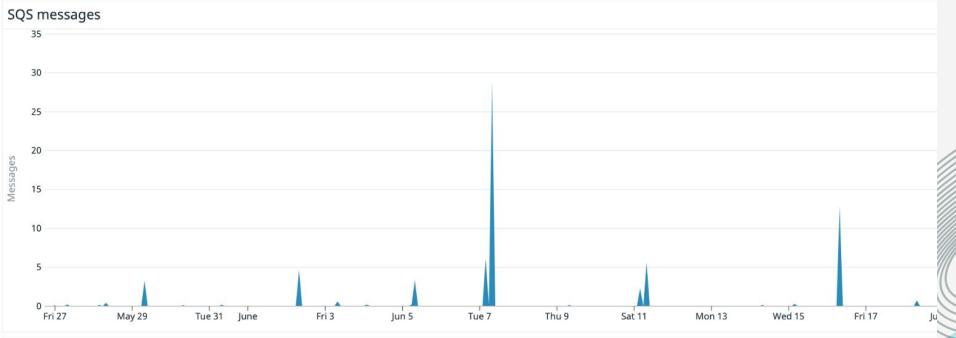
# Analysis

- Stability: crashing after just some hours
- Resiliency: any failure lead to crash the POD
- Performance: slow, opening a new tab in every request
- Scalability: did not handle peaks, static POD number
- High cost: needed to oversize the service



# Figure out what's going on: instrumentation

- Adding and tagging logs
- Measures: AWS built-in, custom
- Know your service!



# Solutions found

- Stability: check the code to release resources, close connections, etc.
- Resiliency: error handlers to catch unexpected errors
- Performance: pool of tabs up & running to handle requests
- Scalability: HPA setup & adding a circuit breaker
- Cost-specific improvements:
  - Chromium fine tuning
  - Screenshot fine tuning
  - AWS spot instances (much cheaper than regular instances)



**DEMO**



**BRACE YOURSELVES**



# Conclusions

- AWS lambda is powerful...
- ...but it is not a swiss knife
- Nice for developing quickly: PoC, MVPs
- If you are running high-intensive workloads  
consider OpenFaaS (or simply k8s services)



Thank  
you!

# David Torres Garrigós



dt@thehotelsnetwork.com



@datoga



<https://github.com/datoga>



<https://linkedin.com/datoga>





[www.thehotelsnetwork.com](http://www.thehotelsnetwork.com)

Headquarters in Barcelona with a team around the world

---

Athens · Austin · Bangkok · Barcelona · Bogotá · Buenos Aires · Frankfurt · Hoi An · Hong Kong · Istanbul  
Jakarta · Los Angeles · Manila · Mexico City · Miami · New York · Paris · San Francisco · Singapore · Vancouver